

```
In [16]: import numpy as np
import skimage
from skimage import transform
from skimage import io
import cv2
import matplotlib.pyplot as plt
from scipy.spatial import distance
import scipy
import random
from PIL import Image
```

```
In [17]: def get_sift(img):
    """
    get the keypoints and compute their SIFT descriptors
    """
    sift = cv2.xfeatures2d.SIFT_create()
    kp, distance = sift.detectAndCompute(img, None)
    return kp, distance

def plot_inlier_matches(ax, img1, img2, inliers):
    """
    plot the match between two image according to the matched keypoints
    """
    res = np.hstack([img1, img2])
    ax.set_aspect('equal')
    ax.imshow(res, cmap='gray')

    ax.plot(inliers[:,0], inliers[:,1], '+r')
    ax.plot(inliers[:,2] + img1.shape[1], inliers[:,3], '+r')
    ax.plot([inliers[:,0], inliers[:,2] + img1.shape[1]],
            [inliers[:,1], inliers[:,3]], 'r', linewidth=0.4)
    ax.axis('off')
```

```
In [18]: #####
### Provided code - nothing to change here ###
#####

"""
Harris Corner Detector
Usage: Call the function harris(filename) for corner detection
Reference (Code adapted from):
    http://www.kaij.org/blog/?p=89
    Kai Jiang - Harris Corner Detector in Python

"""
# Usage:
#harris('./path/to/image.jpg')

def harris(im, min_distance = 10, threshold = 0.1):
    """
```

```

filename: Path of image file
threshold: (optional)Threshold for corner detection
min_distance : (optional)Minimum number of pixels separating
               corners and image boundary
returns: filtered_coords list. index zero is row and index 1 is columns.
"""

harrisim = compute_harris_response(im)
filtered_coords = get_harris_points(harrisim,min_distance, threshold)
plot_harris_points(im, filtered_coords)
return filtered_coords

def gauss_derivative_kernels(size, sizey=None):
    """ returns x and y derivatives of a 2D
        gauss kernel array for convolutions """
    size = int(size)
    if not sizey:
        sizey = size
    else:
        sizey = int(sizey)
    y, x = mgrid[-size:size+1, -sizey:sizey+1]
    #x and y derivatives of a 2D gaussian with standard dev half of size
    # (ignore scale factor)
    gx = - x * exp(-(x**2/float((0.5*size)**2)+y**2/float((0.5*sizey)**2)))
    gy = - y * exp(-(x**2/float((0.5*size)**2)+y**2/float((0.5*sizey)**2)))
    return gx,gy

def gauss_kernel(size, sizey = None):
    """ Returns a normalized 2D gauss kernel array for convolutions """
    size = int(size)
    if not sizey:
        sizey = size
    else:
        sizey = int(sizey)
    x, y = mgrid[-size:size+1, -sizey:sizey+1]
    g = exp(-(x**2/float(size)+y**2/float(sizey)))
    return g / g.sum()

def compute_harris_response(im):
    """ compute the Harris corner detector response function
        for each pixel in the image"""
    #derivatives
    gx,gy = gauss_derivative_kernels(3)
    imx = signal.convolve(im,gx, mode='same')
    imy = signal.convolve(im,gy, mode='same')
    #kernel for blurring
    gauss = gauss_kernel(3)
    #compute components of the structure tensor
    Wxx = signal.convolve(imx*imx,gauss, mode='same')
    Wxy = signal.convolve(imx*imy,gauss, mode='same')
    Wyy = signal.convolve(imy*imy,gauss, mode='same')
    #determinant and trace
    Wdet = Wxx*Wyy - Wxy**2
    Wtr = Wxx + Wyy
    return Wdet / Wtr

def get_harris_points(harrisim, min_distance=10, threshold=0.1):

```

```

""" return corners from a Harris response image
    min_distance is the minimum nbr of pixels separating
    corners and image boundary"""
#find top corner candidates above a threshold
corner_threshold = max(harrisim.ravel()) * threshold
harrisim_t = (harrisim > corner_threshold) * 1
#get coordinates of candidates
candidates = harrisim_t.nonzero()
coords = [ (candidates[0][c],candidates[1][c]) for c in range(len(candidates[0])
#...and their values
candidate_values = [harrisim[c[0]][c[1]] for c in coords]
#sort candidates
index = argsort(candidate_values)
#store allowed point locations in array
allowed_locations = zeros(harrisim.shape)
allowed_locations[min_distance:-min_distance,min_distance:-min_distance] = 1
#select the best points taking min_distance into account
filtered_coords = []
for i in index:
    if allowed_locations[coords[i][0]][coords[i][1]] == 1:
        filtered_coords.append(coords[i])
        allowed_locations[(coords[i][0]-min_distance):(coords[i][0]+min_distance)
        (coords[i][1]-min_distance):(coords[i][1]+min_distance)] = 0
return filtered_coords

def plot_harris_points(image, filtered_coords):
    """ plots corners found in image"""
    figure()
    gray()
    imshow(image)
    plot([p[1] for p in filtered_coords],[p[0] for p in filtered_coords],'r*')
    axis('off')
    show()

#####
### Provided code end                                     ###
#####

```

Your implementations

```

In [19]: def get_matches(img1, img2, percentile):
    kpoint1, dis1 = get_sift(img1)
    kpoint2, dis2 = get_sift(img2)
    kpoint1, kpoint2 = np.array(kpoint1), np.array(kpoint2)

    dist = scipy.spatial.distance.cdist(dis1, dis2, 'sqeuclidean')

    # Flatten the matrix and find the N-th percentile value
    threshold = np.percentile(dist, percentile)
    print(f"threshold is = {threshold}")

    # Get the indices of elements below the percentile
    indices = np.argwhere(dist < threshold)
    print(f"num of matches allowed {indices.shape}")

```

```

# get putative under distance distances
matches = []
new_matches = []
h,w = dist.shape
for i in range(h):
    for j in range(w):
        if dist[i][j] <= 10000:
            new_matches = list(kpoint1[i].pt+kpoint2[j].pt)
            matches.append(new_matches)
matches = np.array(matches)

print(f"num matches {matches.shape}")
return matches

def get_residual(H, data):
    count = len(data)
    point1 = np.hstack((data[:,2:], np.ones((count,1))))
    point2 = data[:,2:]

    est = np.zeros((count,2))
    for i in range(count):
        # x, y = x/w, y/w
        tmp = np.dot(H, point1[i])/np.dot(H, point1[i])[-1]
        est[i] = tmp[:2]

    residual = np.linalg.norm(point2 - est, axis=1) ** 2

    return residual

def ransac(data, t, req_inliers, iterations=1000):
    best_H = None
    best_inliers = None
    max_inliers = 0
    best_residual = float('inf')

    for i in range(iterations):
        sub_index = np.random.choice(len(data), 4)
        subset = data[sub_index]

        #get homography
        A = []
        for i in range(4):
            point1 = np.append(subset[i][:2], 1)
            point2 = np.append(subset[i][2:], 1)

            coord1 = [0,0,0, point1[0], point1[1], point1[2], -point2[1]*point1[0], -point2[2]*point1[0]]
            coord2 = [point1[0], point1[1], point1[2], 0,0,0, -point2[0]*point1[0], -point2[1]*point1[0]]
            A.append(coord1)
            A.append(coord2)

        #SVD
        A = np.array(A)
        U, S, V = np.linalg.svd(A)

```

```

H = V[-1,:].reshape((3,3))

#this is homography transformation
H = H/H[2,2]

if np.linalg.matrix_rank(H) < 3:
    #H is degenerate
    continue

# Get inliers
errors = get_residual(H, data)
index = np.where(errors < t)[0]
inliers = data[index]

# update the model if inliers more than required
if len(inliers) >= req_inliers and len(inliers) >= max_inliers:
    curr_err = errors[index].sum() / len(inliers)
    if len(inliers) == max_inliers:
        # check the residual
        if curr_err > best_residual:
            continue

    best_residual = errors[index].sum() / len(inliers)
    best_inliers = inliers.copy()
    best_H = H.copy()
    max_inliers = len(inliers)

return best_H, best_inliers, max_inliers, best_residual

```

```

In [20]: import numpy as np
from skimage.transform import ProjectiveTransform, warp
from skimage import io

def warp_images(img1, img2, H):
    """
    Warp the right image onto the left image using the provided homography matrix.

    Args:
    - left_image: left color image.
    - right_image: right color image.
    - homography_matrix (numpy.array): 3x3 homography transformation matrix.

    Returns:
    - The composite panorama image.
    """
    transform = ProjectiveTransform(H)

    # Calculate the boundary of warped image
    h, w, z = img1.shape
    raw = np.array([[0, 0], [0, h], [w, h], [w, 0]])
    trans = transform(raw)
    combine = np.vstack((raw, trans))
    min_dot = np.int32(combine.min(axis=0))

```

```

max_dot = np.int32(combine.max(axis=0))
offset_x = -min_dot[0]
offset_y = -min_dot[1]
print("offset is ", (offset_x, offset_y))
T_matrix = np.array([[1, 0, offset_x], [0, 1, offset_y], [0, 0, 1]])
match = cv2.warpPerspective(img2, T_matrix.dot(H), tuple(max_dot - min_dot))
h1, w1, z1 = match.shape

for y in range(h):
    for x in range(w):
        if offset_x + x < w1 and offset_y + y < h1:
            if np.count_nonzero(match[offset_y + y, offset_x + x]) == 0:
                match[offset_y + y, offset_x + x] = img1[y, x]
            else:
                match[offset_y + y, offset_x + x] = (match[offset_y + y, offset
return match

```

Main functions

In [21]: *# Load images*

```

img_left = np.array(Image.open('pier/pier1.JPG'))
img_center = np.array(Image.open('pier/pier2.JPG'))
img_right = np.array(Image.open('pier/pier3.JPG'))

gray_l = Image.open('pier/pier1.JPG').convert("L")
gray_c = Image.open('pier/pier2.JPG').convert("L")
gray_r = Image.open('pier/pier3.JPG').convert("L")
gray_l = cv2.cvtColor(np.array(gray_l), cv2.COLOR_GRAY2BGR)
gray_c = cv2.cvtColor(np.array(gray_c), cv2.COLOR_GRAY2BGR)
gray_r = cv2.cvtColor(np.array(gray_r), cv2.COLOR_GRAY2BGR)

```

In [22]: *# compute and display the center and right matching*

```

matches = get_matches(gray_r, gray_c, percentile=1)

```

threshold is = 127807.0
num of matches allowed (17068, 2)
num matches (383, 4)

In [23]: *# performn RANSAC to get the homography for center and right*

```

H, ransac_match, best_inliers, best_model_errors = ransac(matches, 100, 15)
print("Average residual:", np.average(best_model_errors))
print("Inliers:", best_inliers)
fig, ax = plt.subplots(figsize=(20,10))
plot_inlier_matches(ax, gray_c, gray_r, ransac_match)

```

Average residual: 6.748763206930363
Inliers: 309



In [24]: `H.shape`

Out[24]: `(3, 3)`

In [25]: `# warp right to center`

```
im_right_side = warp_images(gray_c, gray_r, H)


color_im_right_side = warp_images(img_center, img_right, H)

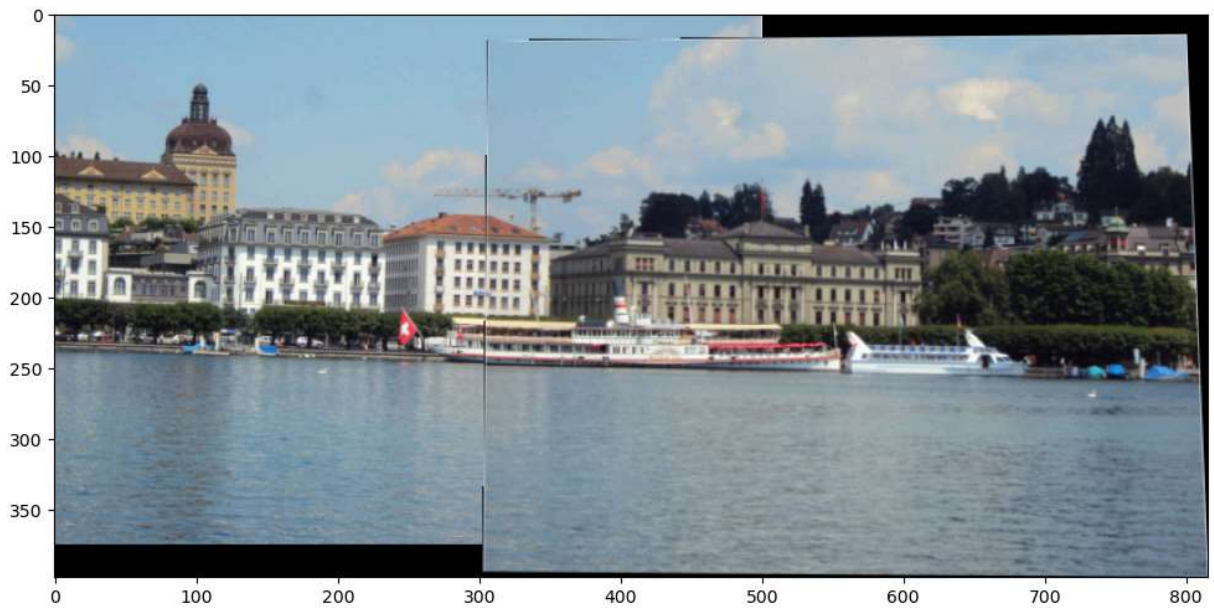
plt.figure(figsize=(12, 8))
plt.imshow(im_right_side)
plt.show()

plt.figure(figsize=(12, 8))
plt.imshow(color_im_right_side)
plt.show()
```

offset is `(0, 0)`

offset is `(0, 0)`





```
In [26]: # compute the left matching with center+ right
matches_l_and_CR = get_matches(gray_l, im_right_side, percentile=10)
```

threshold is = 207780.0
 num of matches allowed (376462, 2)
 num matches (293, 4)

```
In [27]: H, ransac_match, best_inliers, best_model_errors = ransac(matches_l_and_CR, 50, 35,
print("Average residual:", np.average(best_model_errors))
print("Inliers:", best_inliers)
```


Average residual: 1.111529478845029
 Inliers: 276

```
In [28]: im = warp_images(color_im_right_side, img_left, H)
```

```
plt.figure(figsize=(12, 8))
plt.imshow(im)
plt.show()

im = Image.fromarray(im)
im.save("3_image.jpg", format="jpeg")
```

offset is (317, 46)



Extra Credit

```
In [29]: # Apply your stitching code to your own images.
# Load images

img_left = np.array(Image.open('myImages/left.jpg'))
img_center = np.array(Image.open('myImages/center.jpg'))
img_right = np.array(Image.open('myImages/right.jpg'))

gray_l = Image.open('myImages/left.jpg').convert("L")
gray_c = Image.open('myImages/center.jpg').convert("L")
gray_r = Image.open('myImages/right.jpg').convert("L")
gray_l = cv2.cvtColor(np.array(gray_l), cv2.COLOR_GRAY2BGR)
gray_c = cv2.cvtColor(np.array(gray_c), cv2.COLOR_GRAY2BGR)
gray_r = cv2.cvtColor(np.array(gray_r), cv2.COLOR_GRAY2BGR)

# compute and display the center and right matching
matches = get_matches(gray_r, gray_c, percentile=1)

threshold is = 108011.15
num of matches allowed (8865, 2)
num matches (95, 4)

In [30]: # performn RANSAC to get the homography for center and right
```

```
H, ransac_match, best_inliers, best_model_errors = ransac(matches, 100, 15)
print("Average residual:", np.average(best_model_errors))
print("Inliers:", best_inliers)
fig, ax = plt.subplots(figsize=(20,10))
plot_inlier_matches(ax, gray_c, gray_r, ransac_match)
```

Average residual: 17.767663685789575
Inliers: 84



In []: *# warp right to center*

```
im_right_side = warp_images(gray_c, gray_r, H)

color_im_right_side = warp_images(img_center, img_right, H)

plt.figure(figsize=(12, 8))
plt.imshow(im_right_side)
plt.show()

plt.figure(figsize=(12, 8))
plt.imshow(color_im_right_side)
plt.show()

color_im_right_side.save("extra_credit_2images.jpg", format="jpeg")
```

offset is (0, 26)

offset is (0, 26)





```
In [32]: # compute the left matching with center+ right
matches_l_and_CR = get_matches(gray_l, im_right_side, percentile=10)

H, ransac_match, best_inliers, best_model_errors = ransac(matches_l_and_CR, 50, 35,
print("Average residual:", np.average(best_model_errors))
print("Inliers:", best_inliers)
```

threshold is = 198813.0
num of matches allowed (149651, 2)
num matches (93, 4)
Average residual: 9.83039580466553
Inliers: 91

```
In [ ]: im = warp_images(color_im_right_side, img_left, H)
```

```
plt.figure(figsize=(12, 8))
plt.imshow(im)
plt.show()

im = Image.fromarray(im)
im.save("extra_credit_3images.jpg", format="jpeg")
```

offset is (397, 0)

